



# Research on NNVM Compiler

USTC Compiler Team 13

Zhu Yiming, Xie Yunting, Lin Zhiqi

# 01

Part One

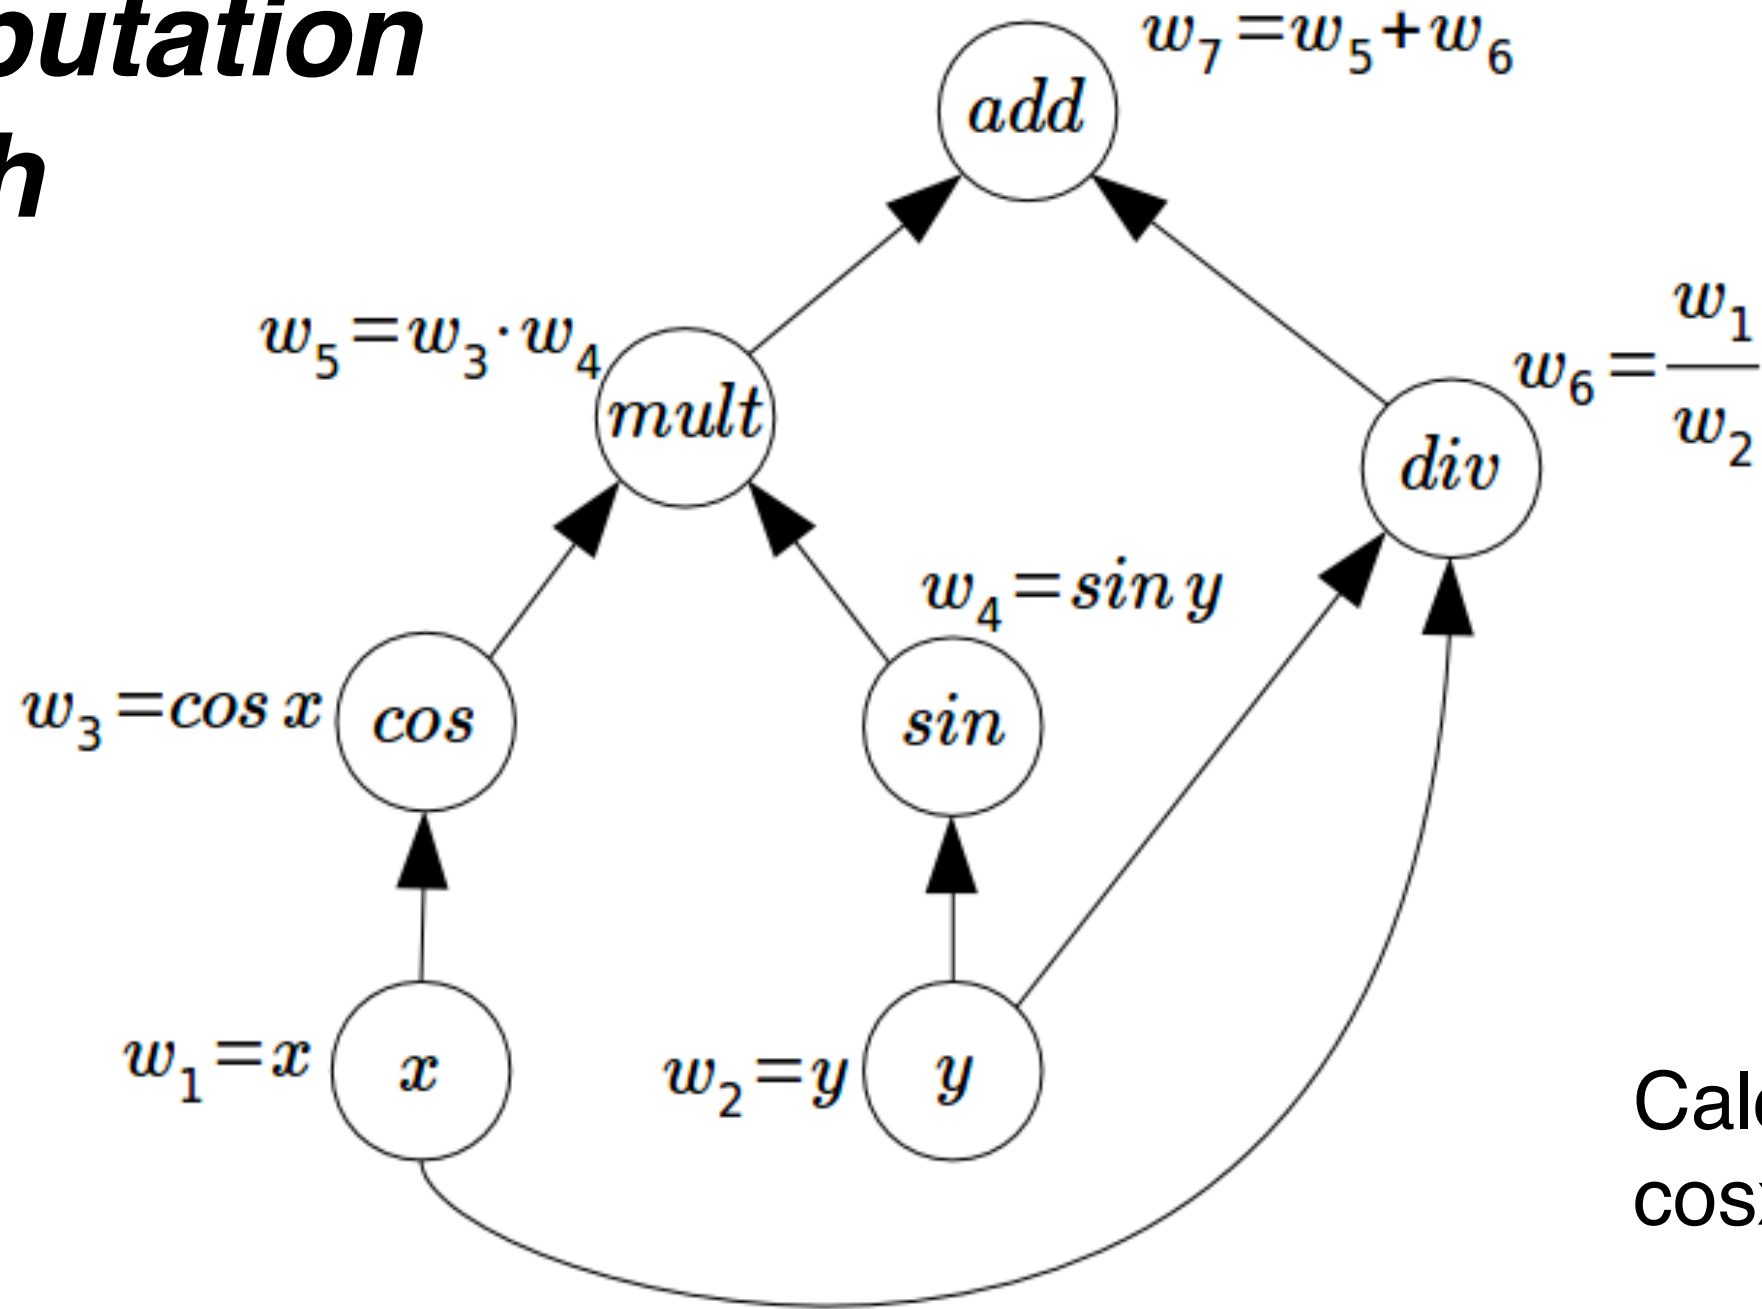
## Background Introduce

NNVM Compiler

nnvm

tvm

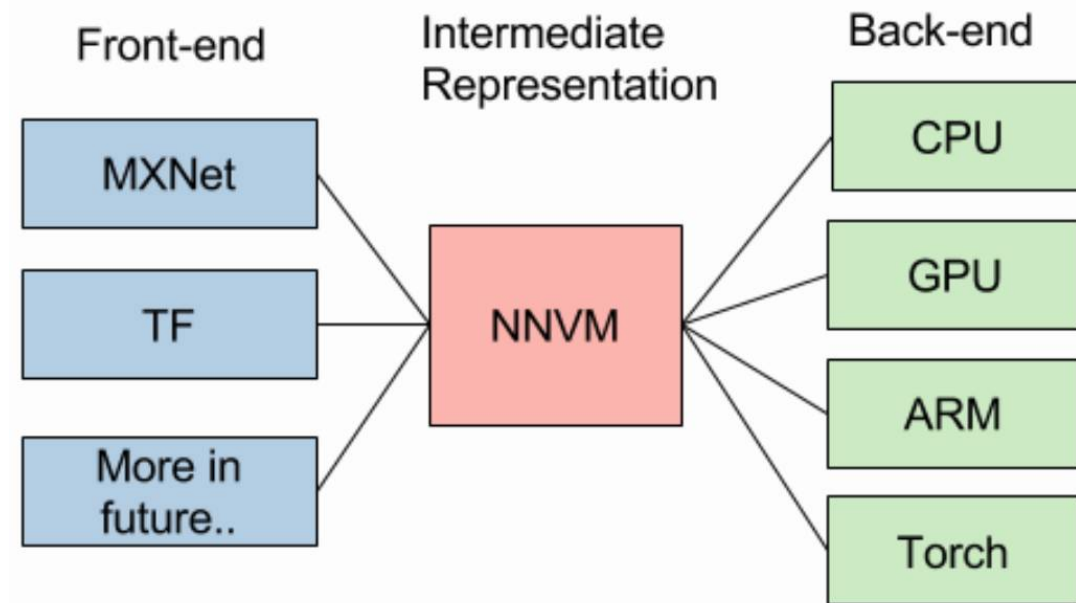
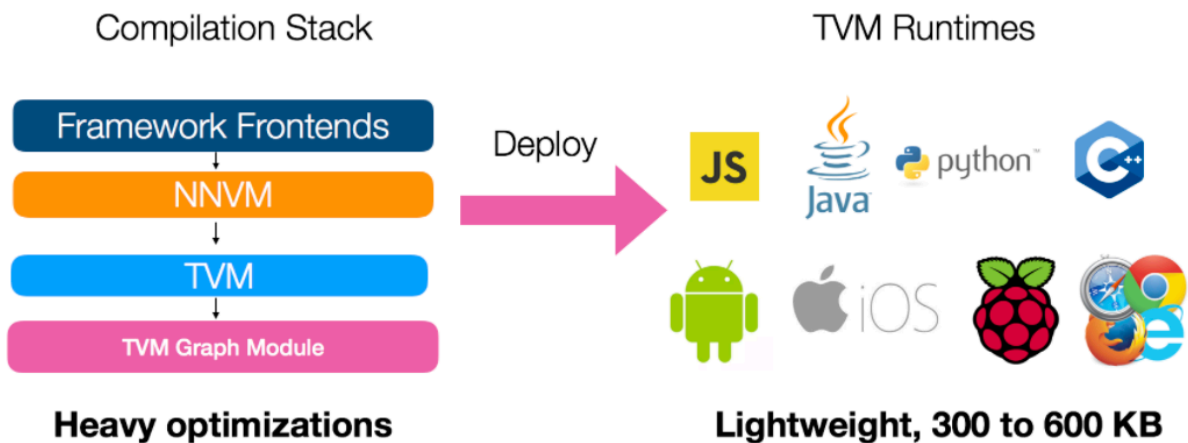
# Computation Graph

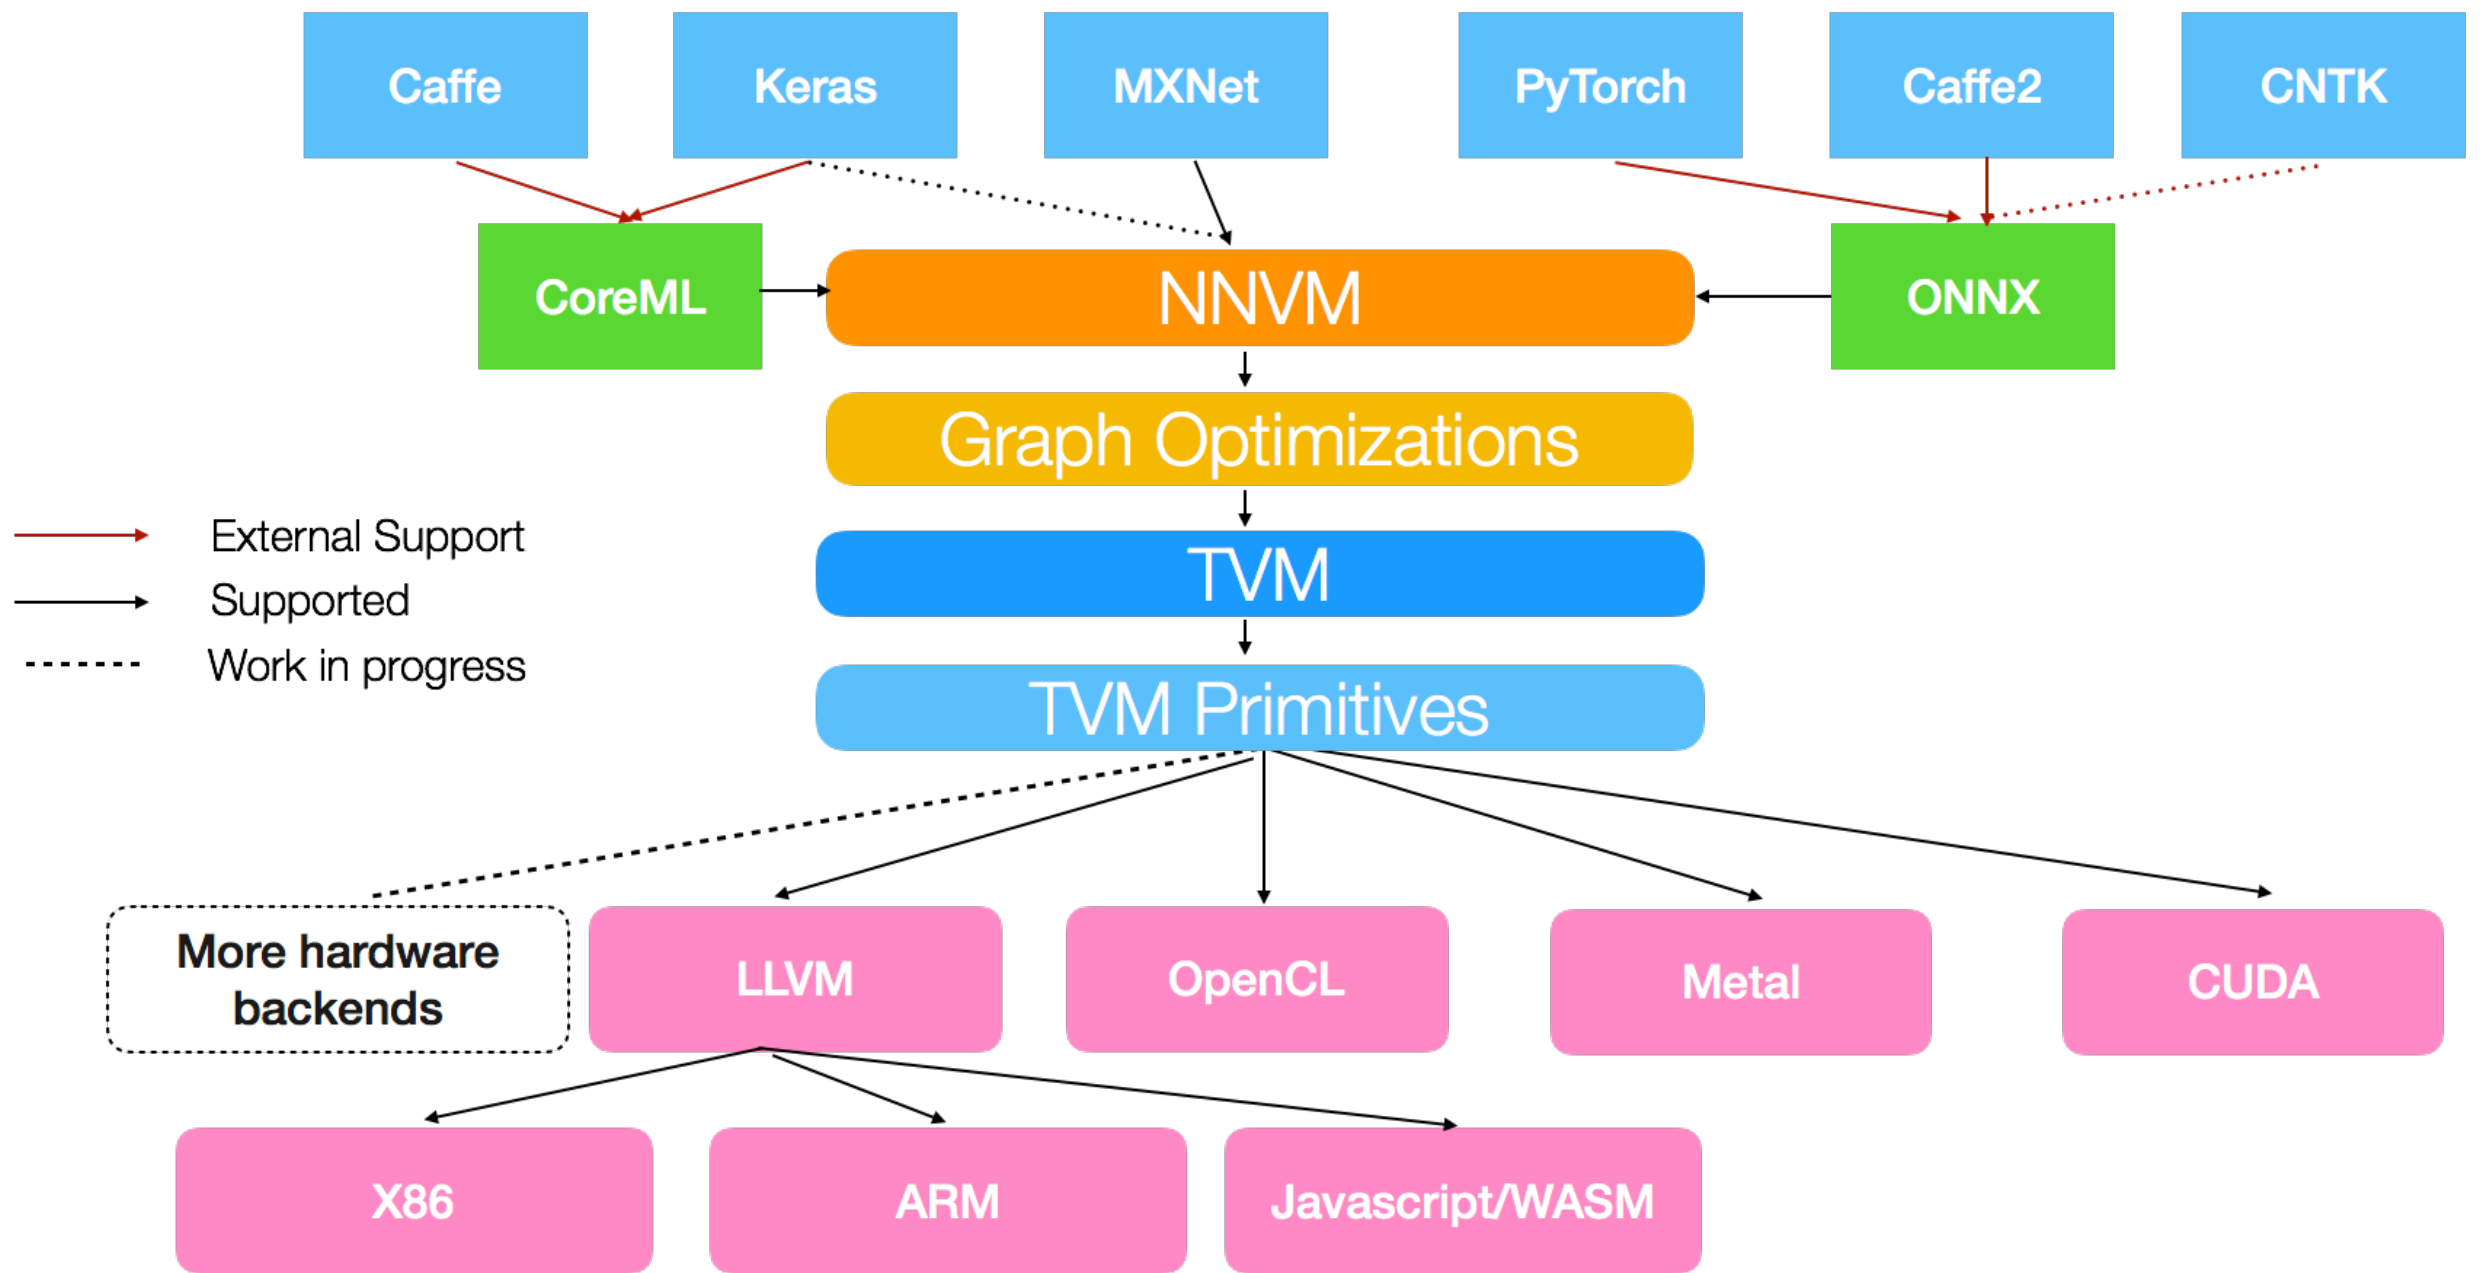




# NNVM Compiler

Neural Network Virtual Machine





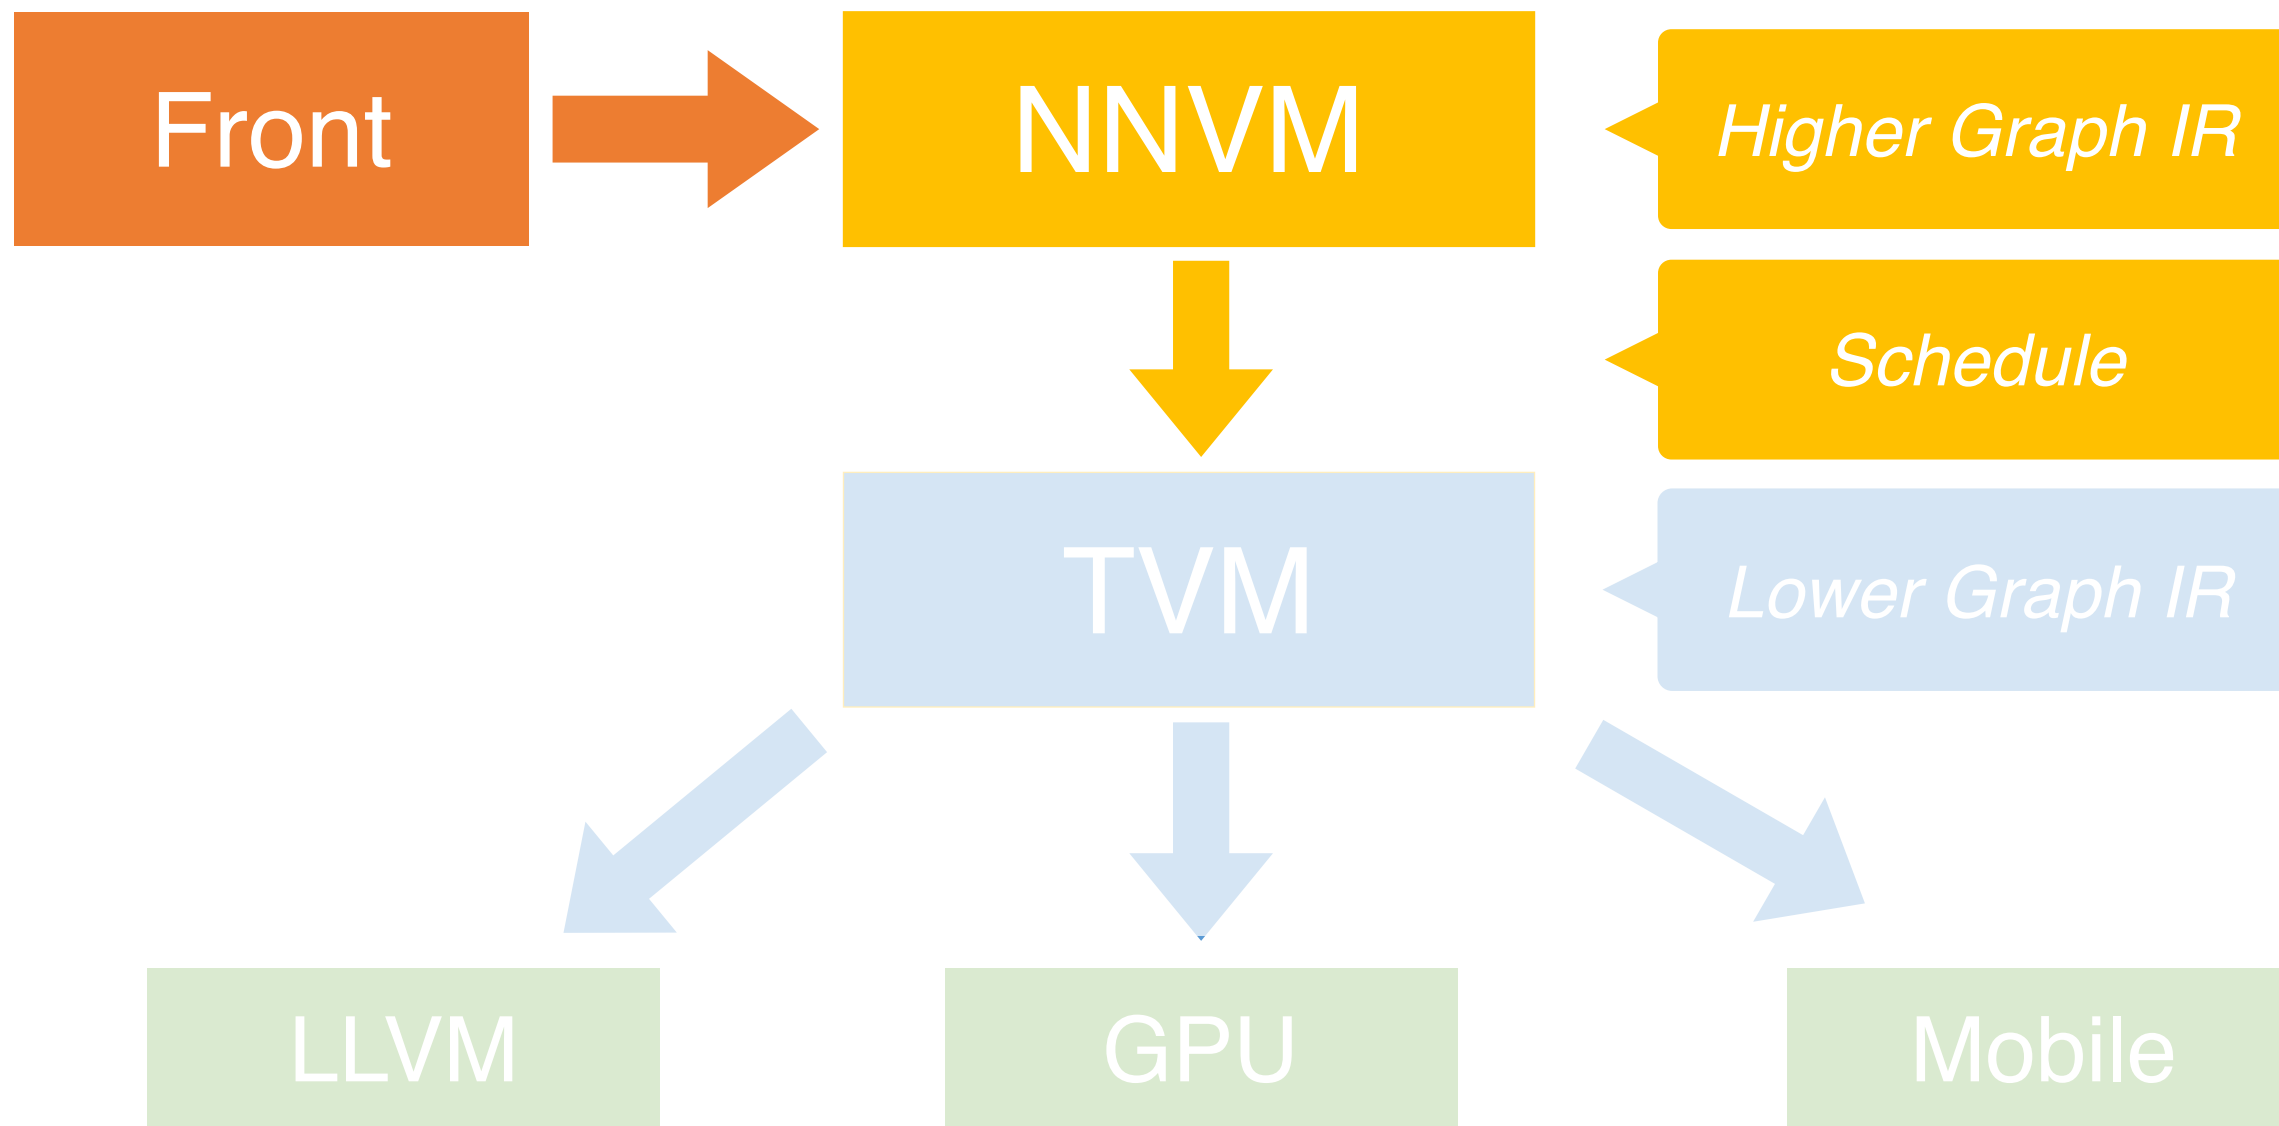
# 02

Part Two

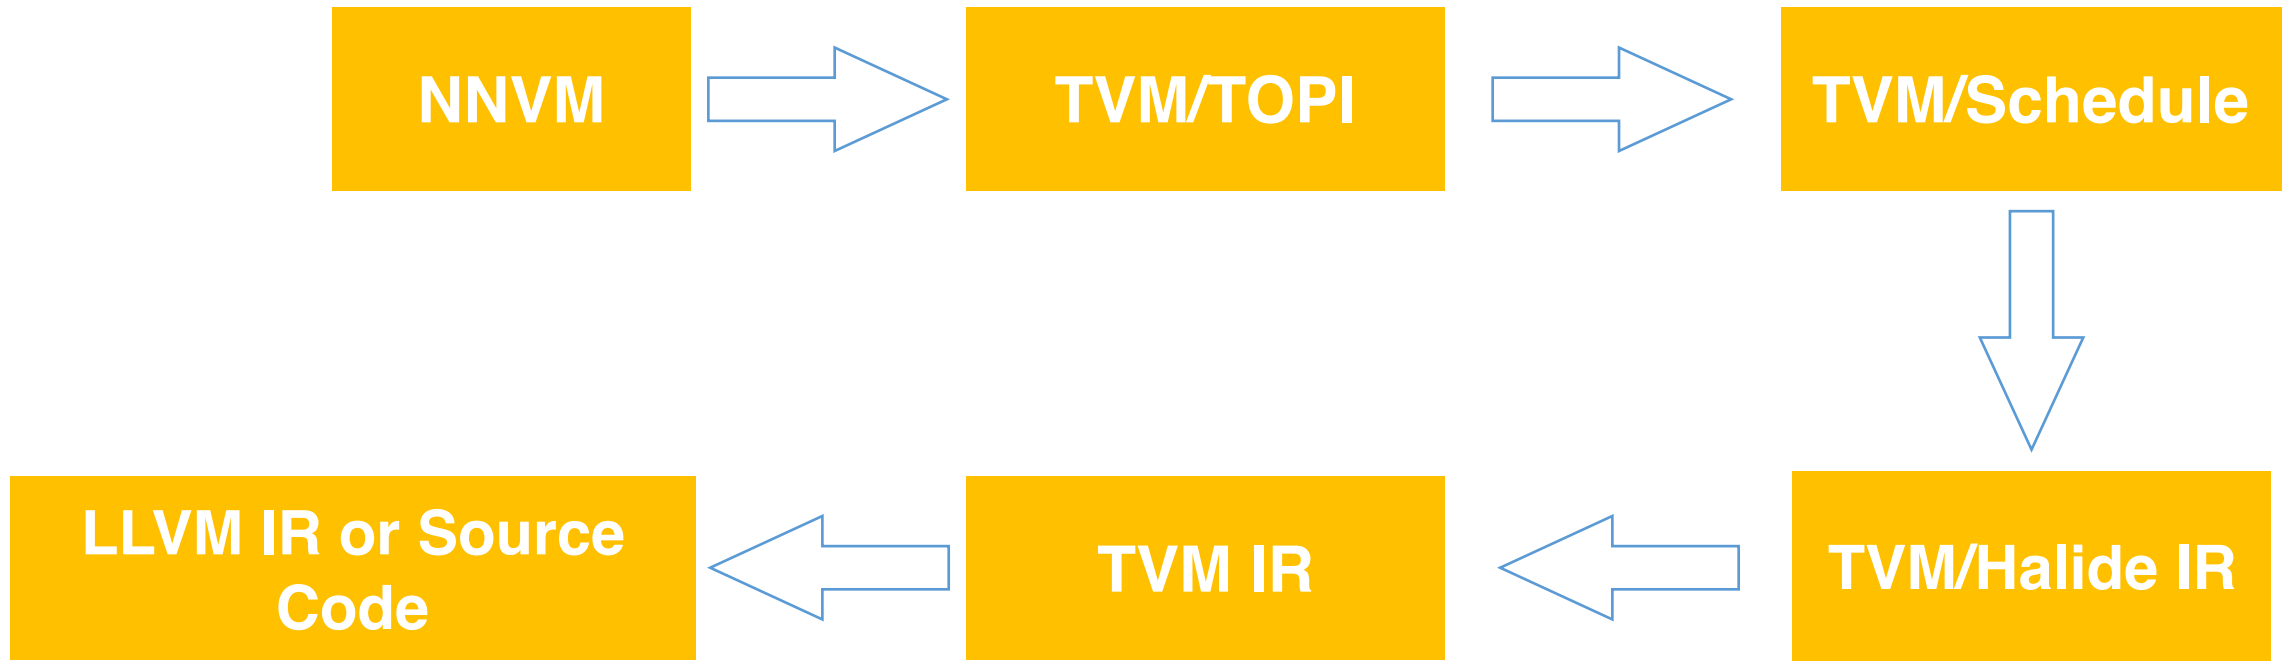
## NNVM Introduction

Design Note  
Optimization

# NNVM Overview



# NNVM Overview



TOPI (TVM Operator Inventory): TOPI is the operator collection library for TVM, to provide sugars for constructing compute declaration as well as optimized schedules.



02



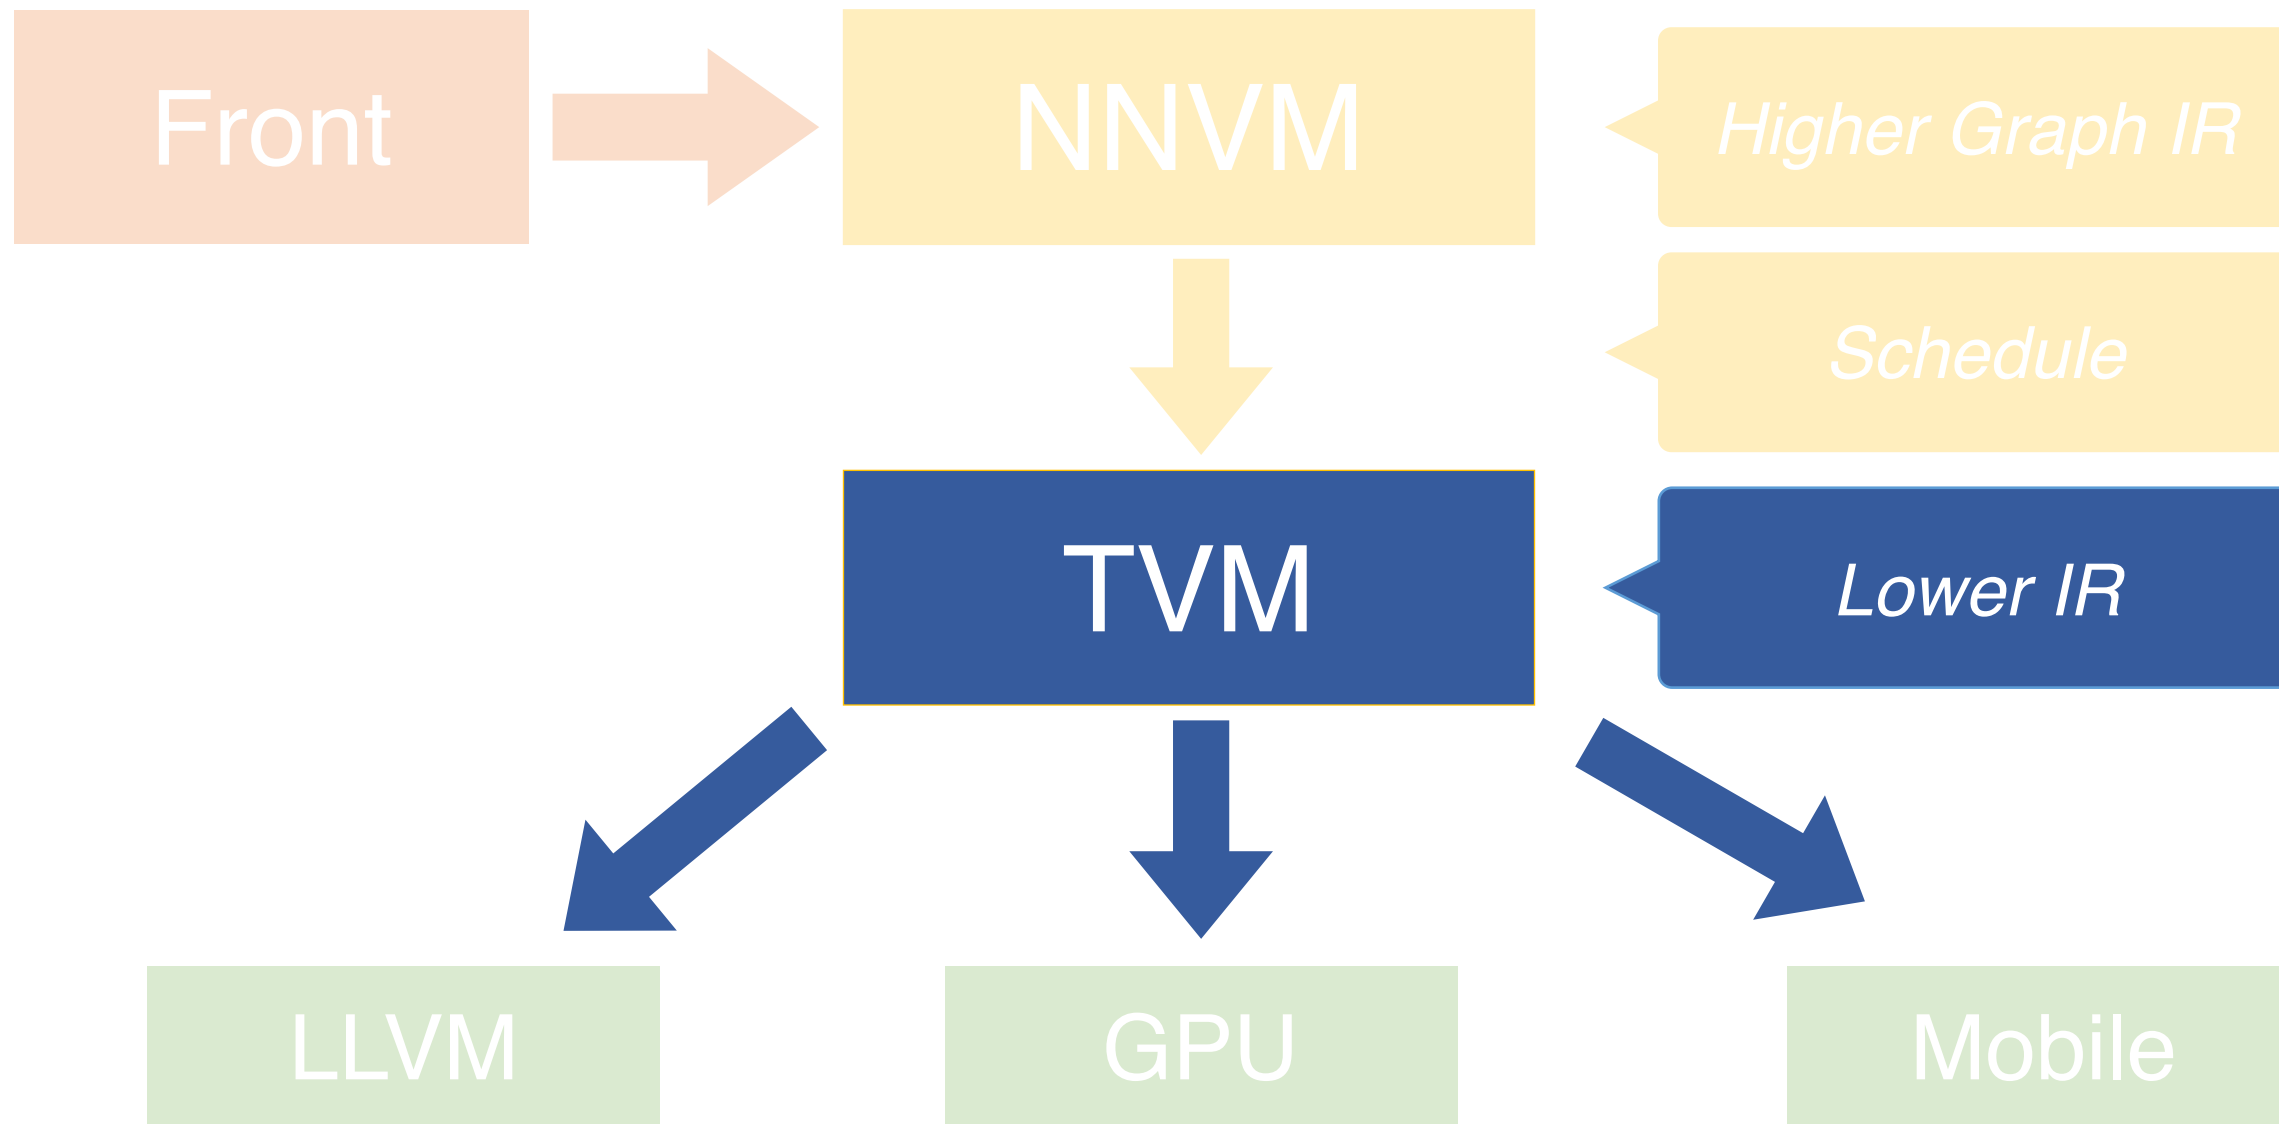
03

Part Three

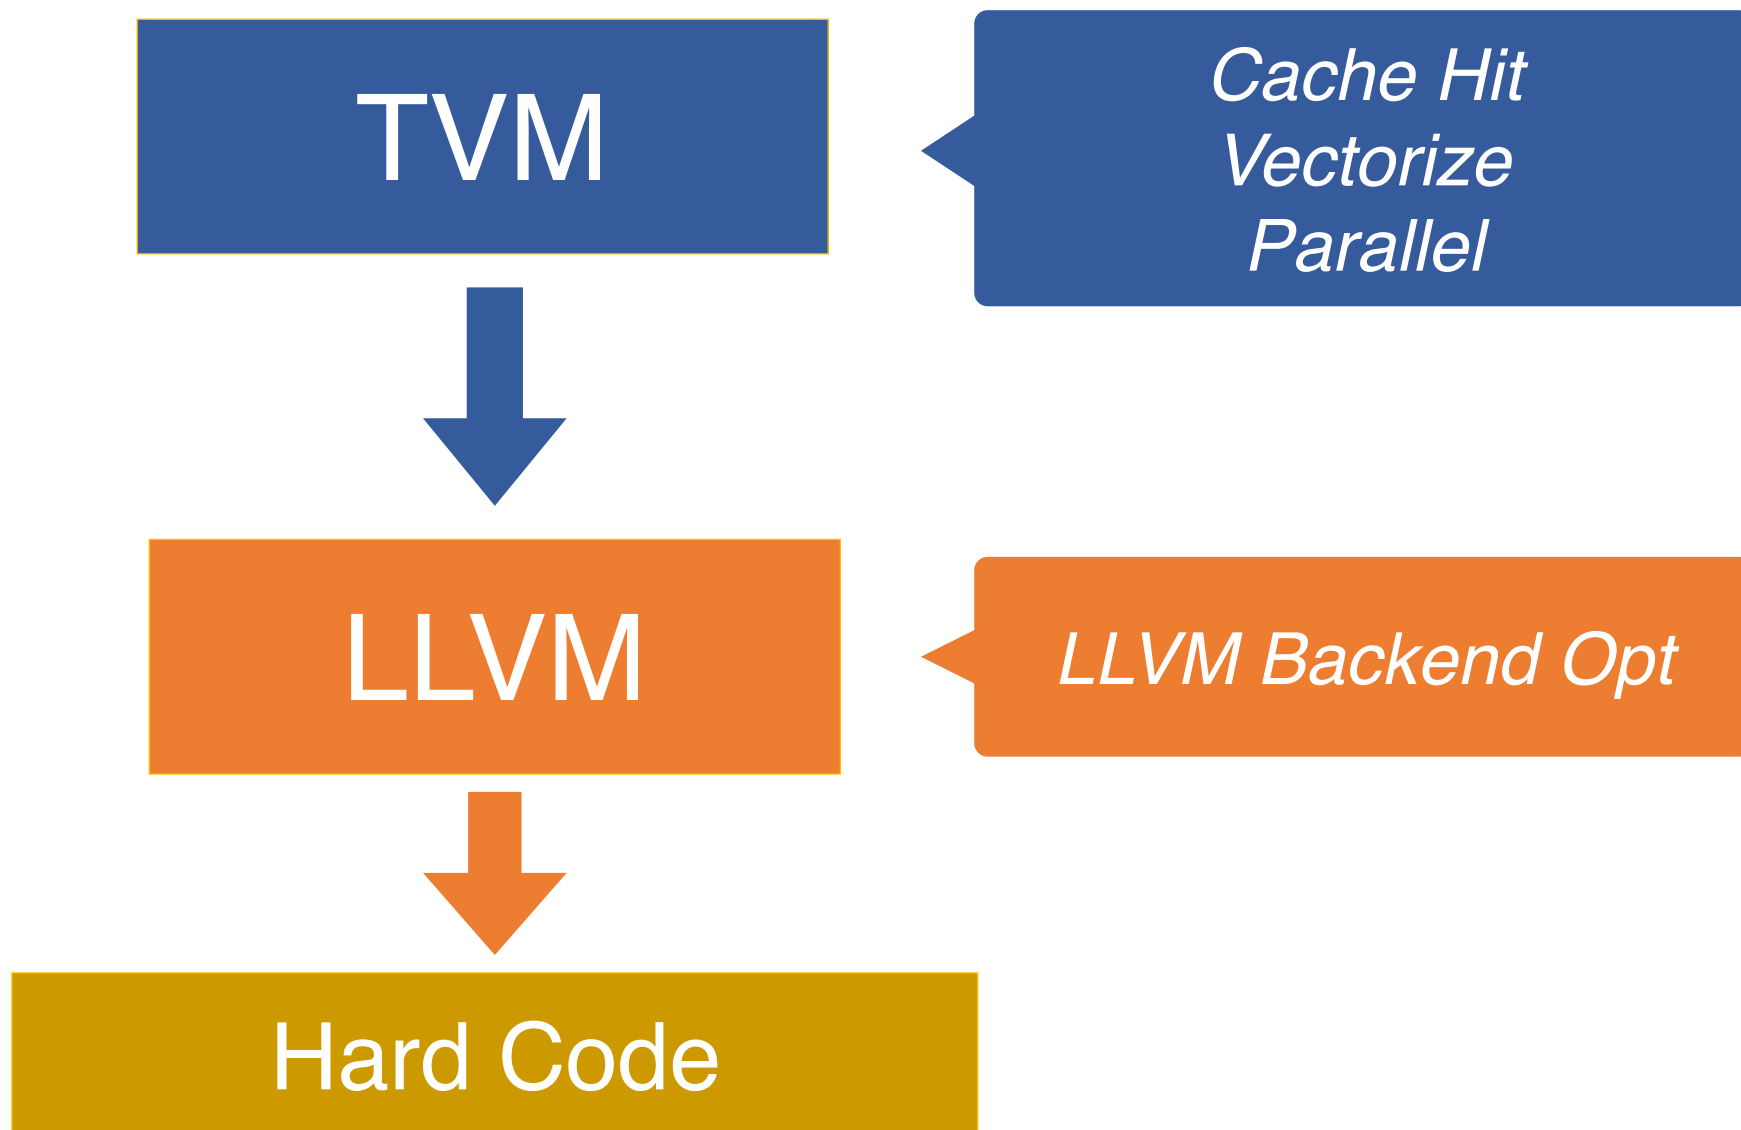
## TVM Stack

Lower IR  
Optimization

# TVM Overview



# TVM Optimization on CPU



# TVM Optimization on CPU - Matrix Multi

```
k = tvm.reduce_axis((0, N), 'k')
A = tvm.placeholder((N, N), name = 'A')
B = tvm.placeholder((N, N), name = 'B')
C = tvm.compute(
    A.shape,
    lambda x, y: tvm.sum(A[x, k] * B[k, y], axis = k),
    name = 'C')
```

*Intel(R) Xeon(R)  
CPU E5-2660  
16 logical cores*

$$C = A \cdot B$$

**T = 2.11167s**


*Matrix dot:*

***A[1024,1024], B[1024,1024]***

```
produce C {
  for (x, 0, 1024) {
    for (y, 0, 1024) {
      C[((x*1024) + y)] = 0.000000f
      for (k, 0, 1024) {
        C[((x*1024) + y)] = (C[((x*1024) + y)] + (A[((x*1024) + k)]*B[(y + (k*1024))]))
      }
    }
  }
}
```

# TVM Optimization on CPU - Tiling/Blocking

*Goal: Increase cache hit rate*




```
produce C {  
  for (x.outer, 0, 32) {  
    for (y.outer, 0, 32) {  
      for (x.inner.init, 0, 32) {  
        for (y.inner.init, 0, 32) {  
          C[(((x.outer*1024) + y.outer) + (x.inner.init*32)) * 32 + y.inner.init] = ...  
        }  
      }  
    }  
  }  
  for (k, 0, 1024) {  
    for (x.inner, 0, 32) {  
      for (y.inner, 0, 32) {  
        C[(((x.outer*1024) + y.outer) + (x.inner*32)) * 32 + y.inner] = ...  
      }  
    }  
  }  
}
```

```
bn = 32  
s = tvm.create_schedule(C.op)  
# Blocking by loop tiling  
xo, yo, xi, yi = s[C].tile(C.op.axis[0], C.op.axis[1], bn, bn)
```

**T = 2.11167s -> 0.897558s**

# TVM Optimization on CPU - Vectorize

## SIMD (Single Instruction Multiple Data)



```
s = tvm.create_schedule(C.op)
xo, yo, xi, yi = s[C].tile(C.op.axis[0], C.op.axis[1], bn, bn)
s[C].reorder(xo, yo, k, xi, yi)

# Vectorization
s[C].vectorize(yi)

func = tvm.build(s, [A, B, C], name = 'mmult')
```

T = 0.897558s -> 0.282522s



# TVM Optimization on CPU - Parallel

Whenever parallelism is possible, just use `parallel()` give a hint to tvvm

```
s = tvvm.create_schedule(C.op)
xo, yo, xi, yi = s[C].tile(C.op.axis[0], C.op.axis[1], bn, bn)
s[C].reorder(xo, yo, xi, k, yi)
# vectorize
s[C].vectorize(yi)

# parallel
s[C].parallel(xo)
```

**T = 0.028522s -> 0.015810s**

Tiling + Vectorize

747%

Tiling +  
Vectorize+  
Parallel

1336%

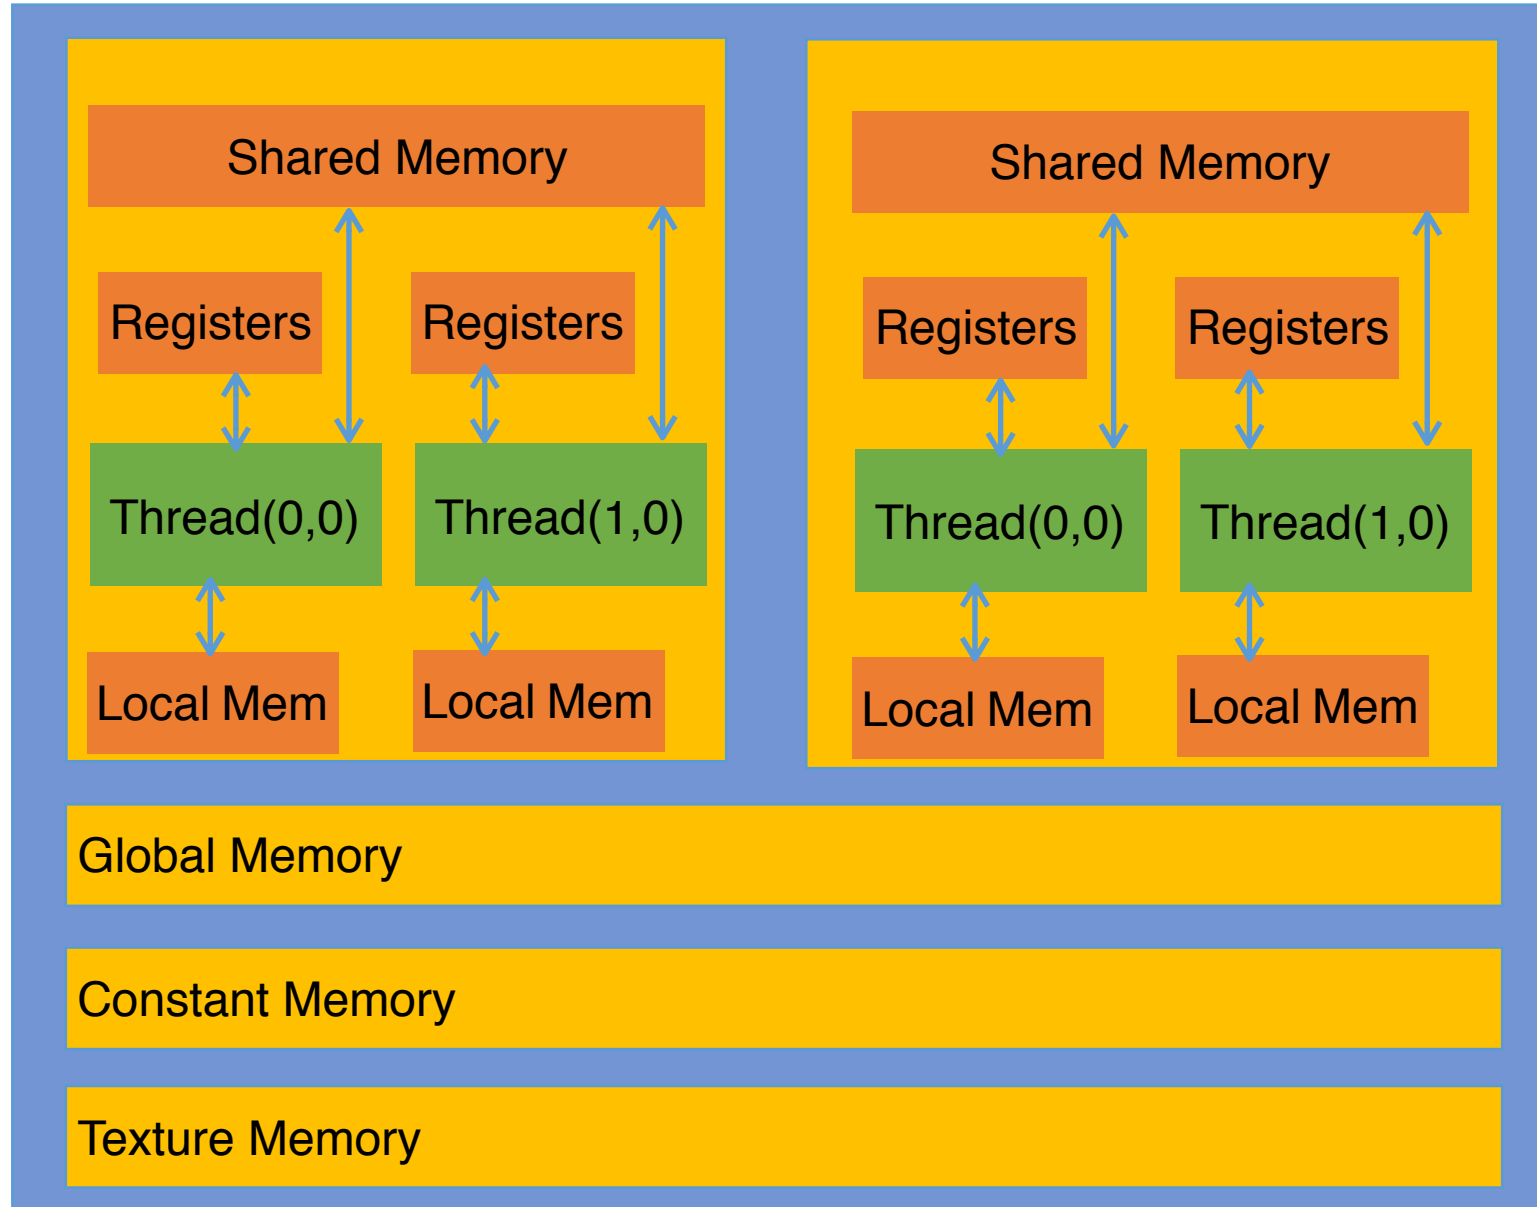
235%

Tiling

# Speedup

Optimization on CPU

# TVM Optimization on GPU



# TVM Optimization on GPU - Convolution

## Convolution

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

\*

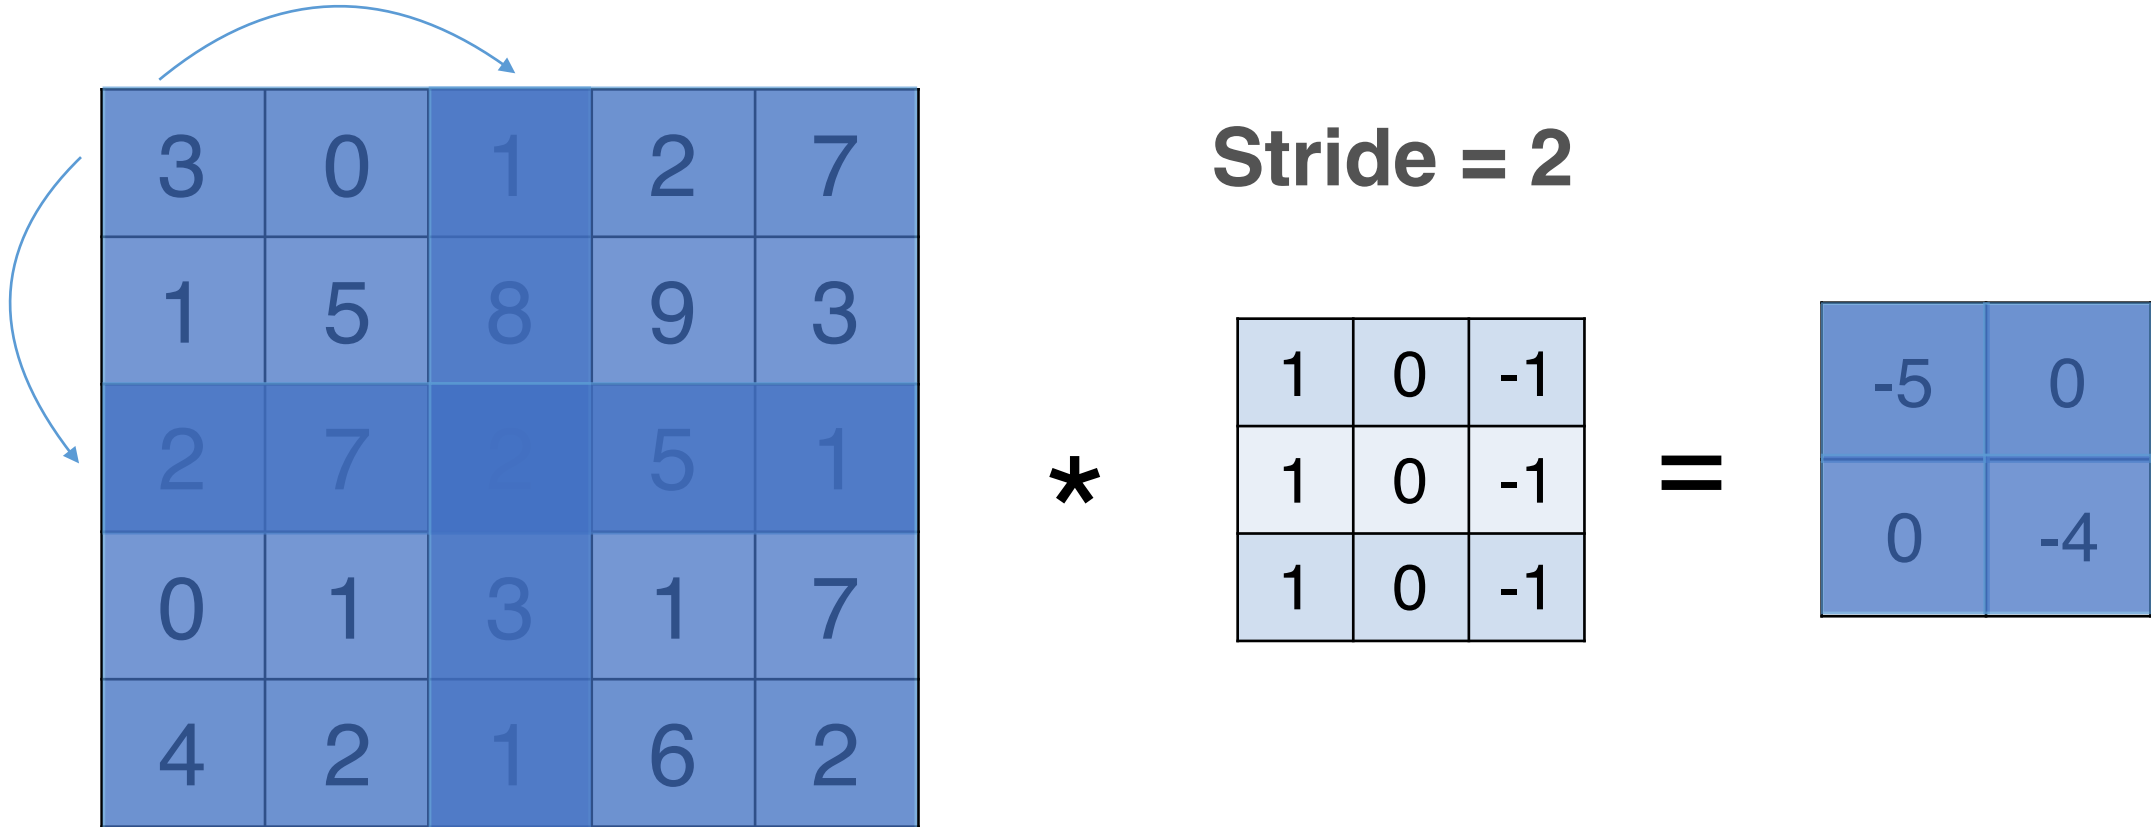
1	0	-1
1	0	-1
1	0	-1

=

-5	-4	0	8
10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

# TVM Optimization on GPU - Convolution

Convolution — Padding — Stride



# TVM Optimization on GPU - Convolution

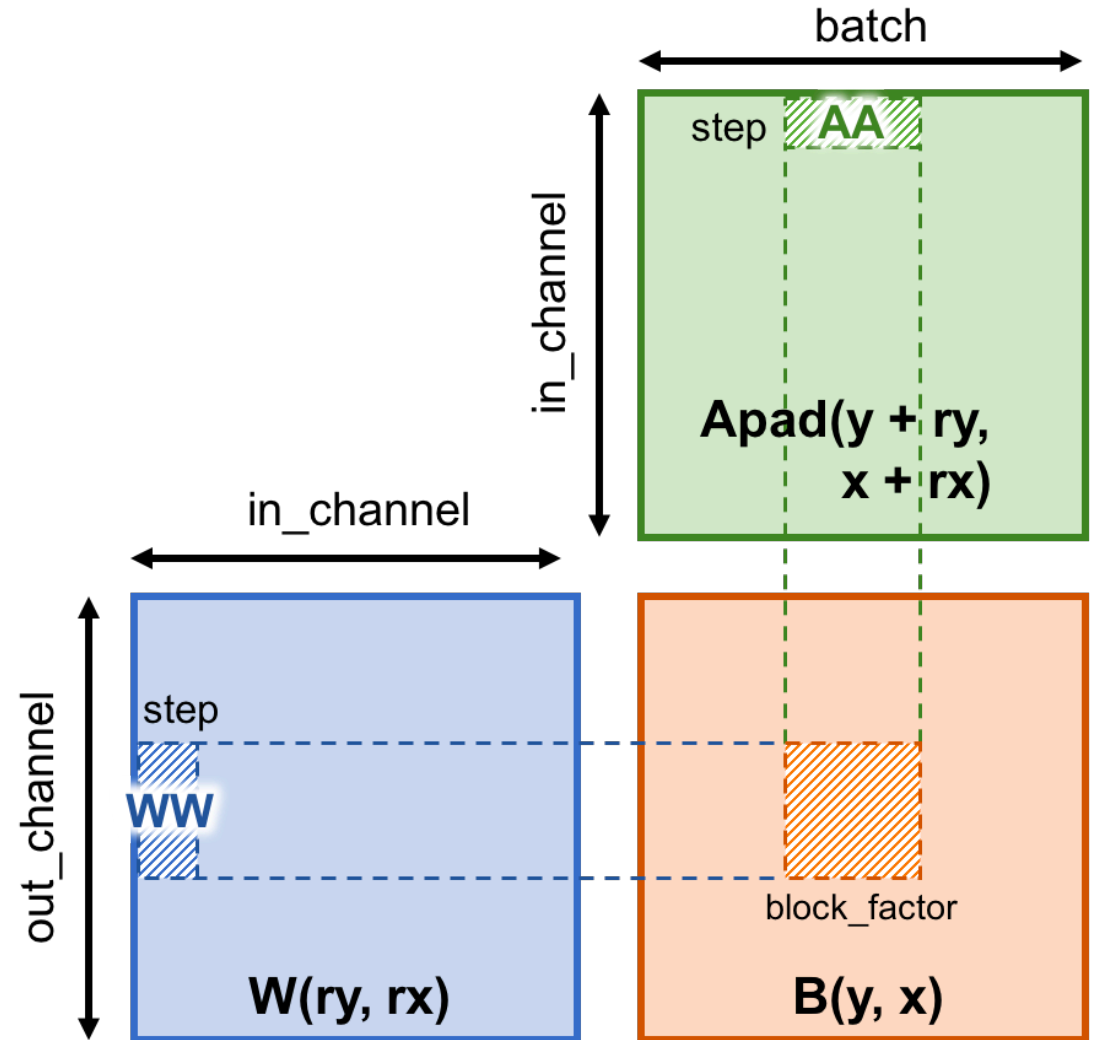
## Optimize1: Blocking

**AA/WW**: Buffer

**Apad**: lhs in convolution

**W**: rhs in convolution

**B**: convolution result






# TVM Optimization on GPU - Convolution

Optimize2:

Virtual Thread Split



Optimize3:

Cooperative Fetching

To reduce memory transfer per thread, make threads in the same thread block cooperatively fetch independent data from global memory

# 04

Part Four

## Compare & Evaluate

XLA  
Darkroom  
Performance

# NNVM vs. XLA

- Different strategy for code generation
- Different strategy for optimization methods
  - XLA: provide a unified optimization method for all hardware resources
  - NNVM: every hardware resource have relevant method
- Similar Optimization on high-level graphs (fusions, layout...)

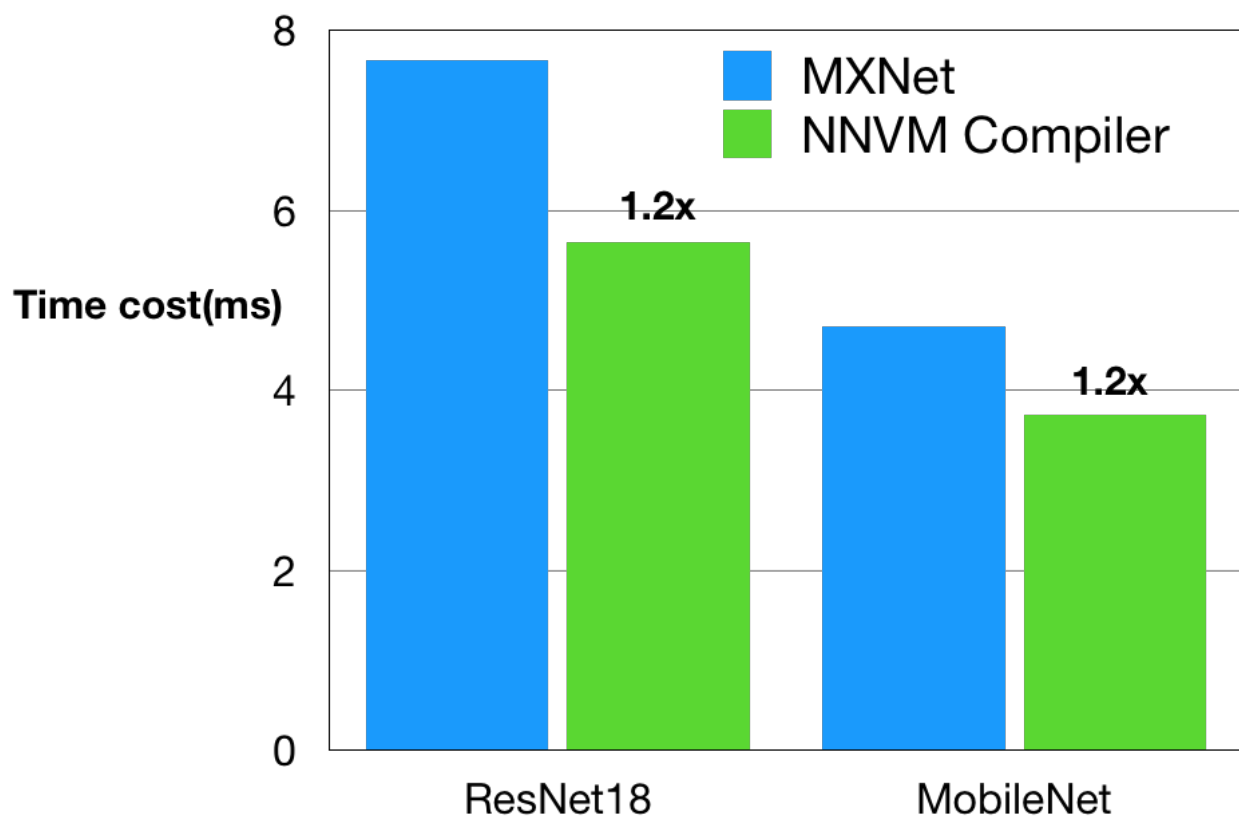
# NNVM vs. Darkroom/Halide

- Similar Ideas between Halide and TVM
  - Both separate algorithms and schedules
- Similar strategy for optimization methods
  - Both have similar optimization on Hardware-independent code
- Different levels of optimization
  - NNVM has more rules and more efficient ways on optimization

# NNVM Performance


## *INVIDIA GPU*

Time cost of Inference on K80



## *Raspberry PI*

Time cost of Inference on Raspberry PI



A photograph of a modern dining area featuring a light-colored wooden table and white plastic chairs with wooden legs. The room has large windows in the background, letting in natural light. A teal-colored geometric overlay is present in the top-left and bottom-right corners of the image.

# Thanks

USTC Compiler Team 13 nnvm-tvm  
Zhu Yiming, Xie Yunting, Lin Zhiqi